

Search


Update a Cookbook to Dynamically Use Nodes with the `company_web` policy name

Objectives

After completing this module, you should be able to

- Describe the query syntax used in search
- Build a search into your recipe code
- Create a Ruby String and Array dynamically
- Test that your load balancer is still balancing traffic

CONCEPT



Search

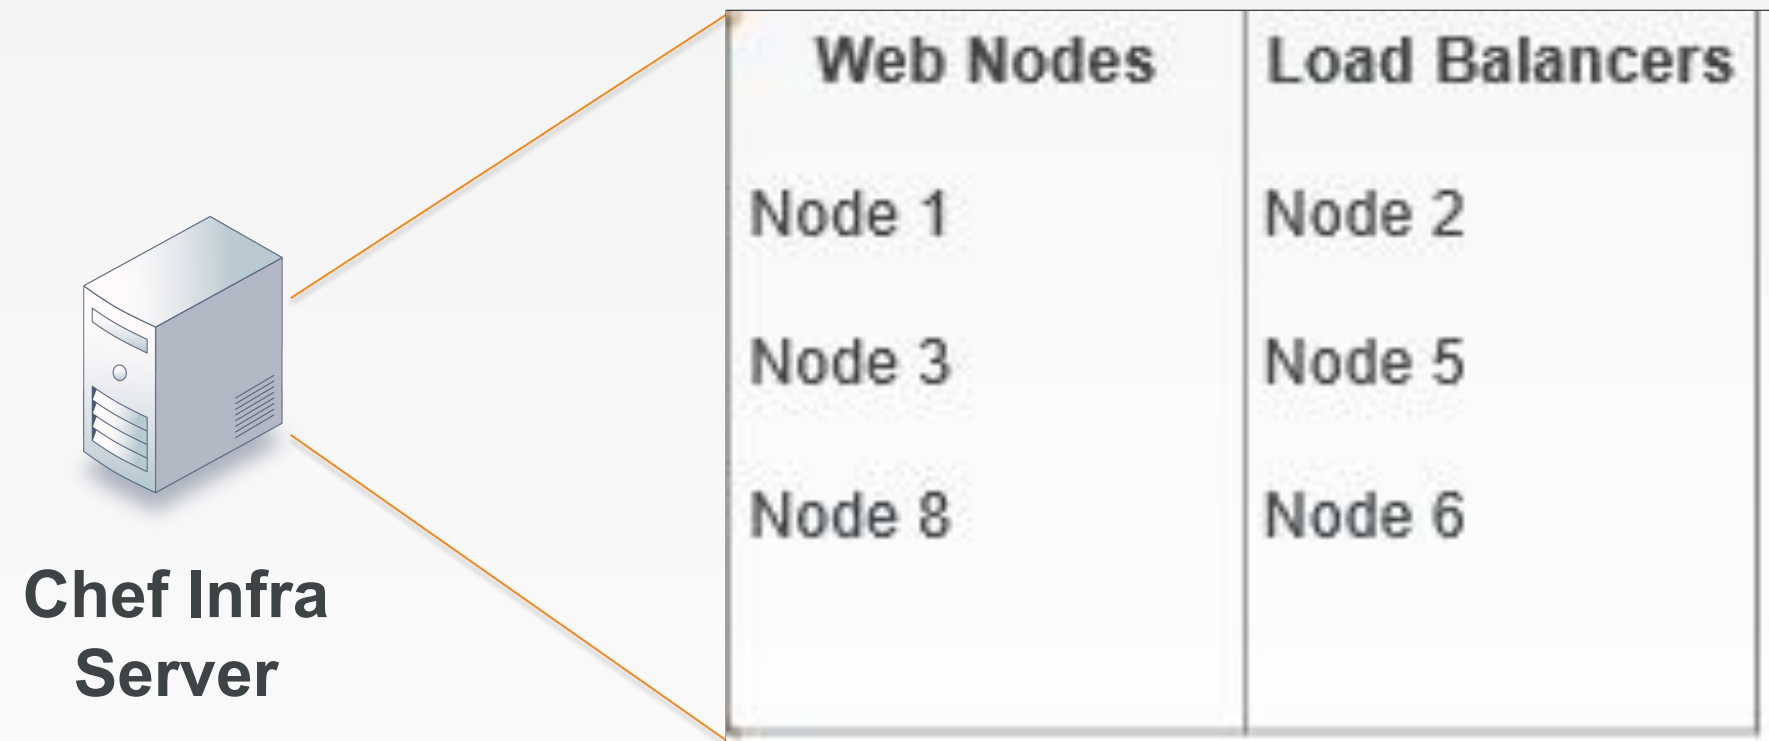
To add new servers as load balancer members, we would need to bootstrap a new web server and then update our load balancer's myhaproxy cookbook recipe.

That seems inefficient to have to update a cookbook recipe manually.

The Chef Infra Server and Search

Chef Infra Server maintains a representation of all the nodes within our infrastructure that can be searched on.

Search is a service discovery tool that allows us to query the Chef Infra Server.

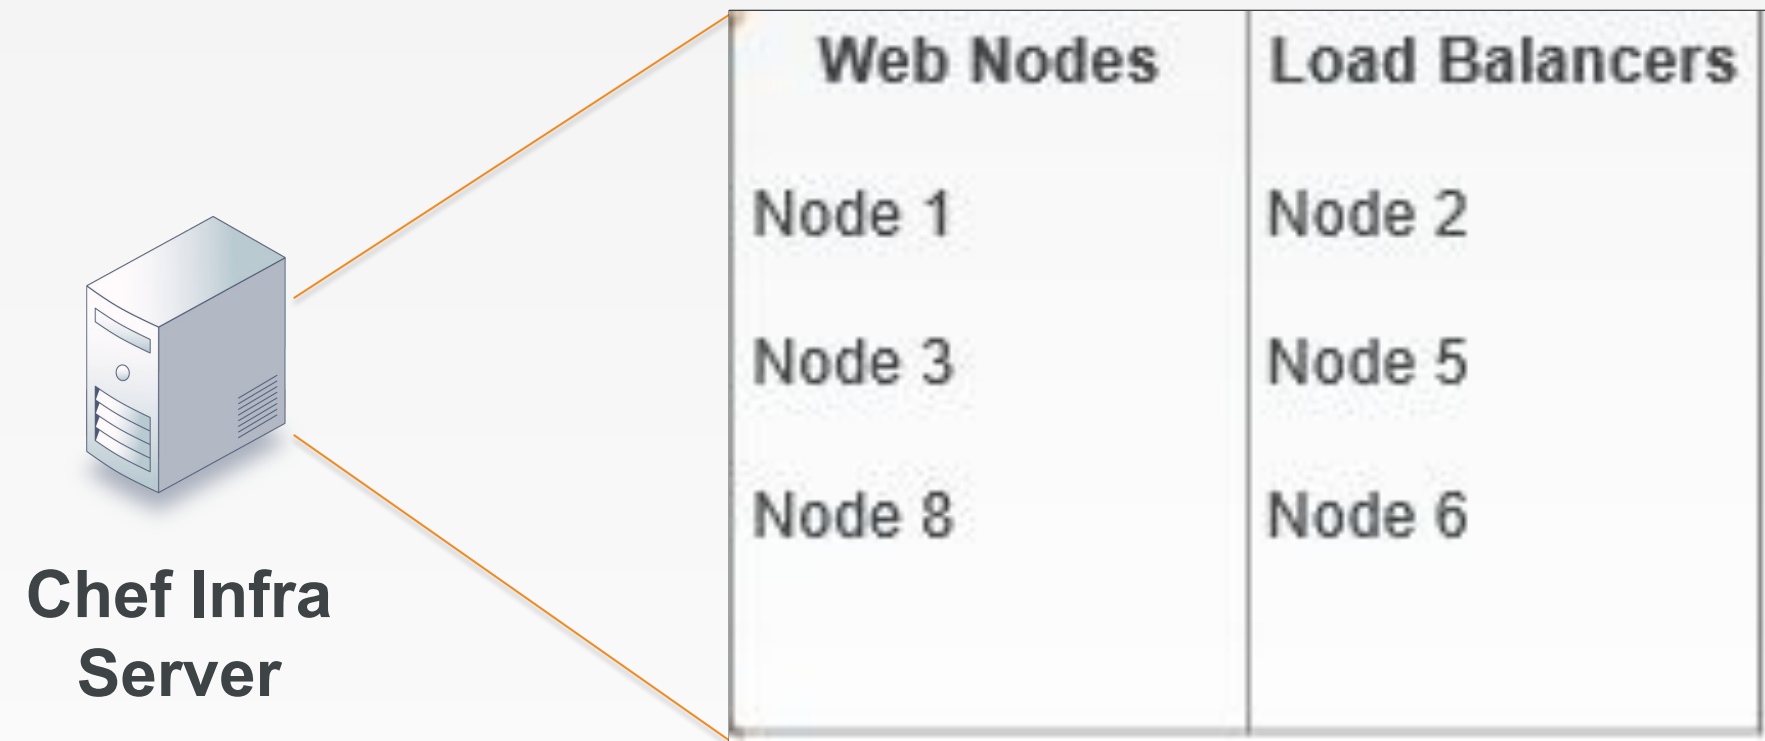


https://docs.chef.io/chef_search.html

https://docs.chef.io/chef_search.html#search-indexes

The Chef Infra Server and Search

We can ask the Chef Infra Server to return all the nodes or a subset of nodes based on the query syntax that we provide it through ``knife search`` or within our recipes through ``search``.



Search Syntax

A search query is comprised of two parts: the key and the search pattern. A search query has the following syntax:

key:search_pattern

...where key is a field name that is found in the JSON description of an indexable object on the Chef Infra Server and search_pattern defines what will be searched for

Search Criteria

We may use wildcards within search so a search criteria that we could use is: "*" : "*"

However, querying and returning every node is not what we need to solve our current problem.

Scenario: We want only to return a subset of our nodes... only the nodes that are web servers.



Demo: View Information for All Nodes



```
> knife search node *:*
```

```
Node Name:  apache_web
Policy Name: company_web
Policy Group: prod
FQDN:       ip-172-31-57-169.ec2.internal
IP:         34.196.104.17
Run List:    recipe[company_web::default]
Recipes:     company_web::default, apache::default, apache::server
Platform:    centos 7.6.1810
Tags:

Node Name:  iis_web
Policy Name: company_web
Policy Group: prod
FQDN:       WIN-DQFQCUFHDCP.ec2.internal
IP:         34.195.38.226
Run List:    recipe[company_web::default]
Recipes:     company_web::default, myiis::default, myiis::server
Platform:    windows 6.3.9600
Tags:

Node Name:  lb
Policy Name: myhaproxy
Policy Group: prod
FQDN:       ip-172-31-22-163.ec2.internal
IP:         34.196.50.77
Run List:    recipe[myhaproxy::default]
Recipes:     myhaproxy::default, haproxy::manual, haproxy::install_package
Platform:    centos 7.6.1810
```


Demo: View Information for All Server Nodes



```
> knife search node policy_name:company_web
```

```
Node Name:    apache_web
Policy Name:   company_web
Policy Group:  prod
FQDN:         ip-172-31-57-169.ec2.internal
IP:           34.196.104.17
Run List:     recipe[company_web::default]
Recipes:      company_web::default, apache::default, apache::server
Platform:     centos 7.6.1810
Tags:
```

```
Node Name:    iis_web
Policy Name:   company_web
Policy Group:  prod
FQDN:         WIN-DQFQCUFHDCP.ec2.internal
IP:           34.195.38.226
Run List:     recipe[company_web::default]
Recipes:      company_web::default, myiis::default, myiis::server
Platform:     windows 6.3.9600
Tags:
```

Demo: Return Public Hostname and IP for Servers



```
> knife search node policy_name:company_web -a cloud
```

```
apache_web:
  cloud:
    local_hostname:    ip-172-31-57-169.ec2.internal
    local_ipv4:        172.31.57.169
    local_ipv4_addrs:  172.31.57.169
    provider:          ec2
    public_hostname:    ec2-34-196-104-17.compute-1.amazonaws.com
    public_ipv4:        34.196.104.17
    public_ipv4_addrs: 34.196.104.17

iis_web:
  cloud:
    local_hostname:    ip-172-31-62-51.ec2.internal
    local_ipv4:        172.31.62.51
    local_ipv4_addrs:  172.31.62.51
    provider:          ec2
    public_hostname:    ec2-34-195-38-226.compute-1.amazonaws.com
    public_ipv4:        34.195.38.226
    public_ipv4_addrs: 34.195.38.226
```

Demo: Return Public Hostname for Servers



```
> knife search node policy_name:company_web -a cloud.public_hostname
```

```
2 items found
```

```
apache_web:
```

```
  cloud.public_hostname: ec2-34-196-104-17.compute-1.amazonaws.com
```

```
iis_web:
```

```
  cloud.public_hostname: ec2-34-195-38-226.compute-1.amazonaws.com
```

Search Syntax within a Recipe

```
web_nodes = search('node', 'policy_name:company_web')
```

creates and names a
variable

assigns the value of the
operation on the right
into the variable on the left

invokes the search method

the index or items to search

the search criteria -
key:value

The search syntax within a recipe differs from the search syntax when using `knife search` from the command line.

Hard Coding Example


```
haproxy_install 'package'

haproxy_frontend 'http-in' do
  bind '*:80'
  default_backend 'server_backend'
end

haproxy_backend 'server_backend' do
  server [
    'ec2-54-175-46-24.compute-1.amazonaws.com 54.175.46.24:80 maxconn 32',
    'ec2-34-196-10-17.compute-1.amazonaws.com 34.196.10.17:80 maxconn 32'
  ]
end

haproxy_service 'haproxy'
```

EXERCISE



GL: Dynamic Web Load Balancer

Every time we create a web node we need to update our load balancer (myhaproxy) cookbook. That doesn't feel right!

Objective:

- ☐ Update the myhaproxy cookbook to dynamically use nodes with the company_web policy_name.
- ☐ Update the major version of the myhaproxy cookbook
- ☐ Upload the Cookbook
- ☐ Run chef-client on the load balancer node
- ☐ Verify the load balancer node relays requests to both web nodes

GL: Remove the Hard-Coded Members

```
~/chef-repo/cookbooks/myhaproxy/recipes/default.rb
```

```
haproxy_install 'package'

haproxy_frontend 'http-in' do
  bind '*:80'
  default_backend 'server_backend'
end
```

```
haproxy_backend 'server_backend' do
  server [
    'ec2-54-175-46-24.compute-1.amazonaws.com 54.175.46.24:80 maxconn 32',
    'ec2-34-196-10-17.compute-1.amazonaws.com 34.196.10.17:80 maxconn 32'
  ]
end
```

```
haproxy_service 'haproxy'
```

GL: Add the Search Criteria



```
~/chef-repo/cookbooks/myhaproxy/recipes/default.rb
```

```
...  
  
haproxy_install 'package'  
  
haproxy_frontend 'http-in' do  
  bind '*:80'  
  default_backend 'server_backend'  
end  
  
web_nodes = search('node', 'policy_name:company_web')
```

Note: We will provide the final recipe in a moment.

GL: Create an Array to Store the Converted Members

`~/chef-repo/cookbooks/myhaproxy/recipes/default.rb`

```
...  
web_nodes = search('node', 'policy_name:company_web')  
  
server_array = []  
  
haproxy_backend 'server_backend' do  
  server server_array  
end  
  
haproxy_service 'haproxy'
```

GL: Create an Array to Store the Converted Members



```
~/chef-repo/cookbooks/myhaproxy/recipes/default.rb
```

```
...
web_nodes = search('node', 'policy_name:company_web')

server_array = []

web_nodes.each do |one_node|
  one_server = ""
  # TODO: Populate the array with each webserver's hostname and ipaddress
end

haproxy_backend 'server_backend' do
  server server_array
end

haproxy_service 'haproxy'
```

GL: Create an Array to Store the Converted Members



```
~/chef-repo/cookbooks/myhaproxy/recipes/default.rb
```

```
...
web_nodes = search('node', 'policy_name:company_web')

server_array = []

web_nodes.each do |one_node|

  one_server = "#{one_node['cloud']['public_hostname']} #{one_node['cloud']['public_ipv4']}:80
maxconn 32"
  server_array.push(one_server)

end

haproxy_backend 'server_backend' do
  server server_array
end

haproxy_service 'haproxy'
```

GL: The Final Recipe



```
~/chef-repo/cookbooks/myhaproxy/recipes/default.rb
```

```
haproxy_install 'package'

haproxy_frontend 'http-in' do
  bind '*:80'
  default_backend 'server_backend'
end

web_nodes = search('node', 'policy_name:company_web')


server_array = []

web_nodes.each do |one_node|
  one_server = "#{one_node['cloud']['public_hostname']} #{one_node['cloud']['public_ipv4']}:80 maxconn
32"
  server_array.push(one_server)
end

haproxy_backend 'server_backend' do
  server server_array
end

haproxy_service 'haproxy'
```

EXERCISE



GL: Dynamic Web Load Balancer

Every time we create a web node we need to update our load balancer (myhaproxy) cookbook. That doesn't feel right!

Objective:


- ✓ Update the myhaproxy cookbook to dynamically use nodes with the company_web policy_name.
- ☐ Update the major version of the myhaproxy cookbook
- ☐ Upload the Cookbook
- ☐ Run chef-client on the load balancer node
- ☐ Verify the load balancer node relays requests to both web nodes

GL: Updating the Cookbook's Version Number

```
~/chef-repo/cookbooks/myhaproxy/metadata.rb
```

```
name 'myhaproxy'  
maintainer 'The Authors'  
maintainer_email 'you@example.com'  
license 'All Rights Reserved'  
description 'Installs/Configures myhaproxy'  
long_description 'Installs/Configures myhaproxy'  
version '1.0.0'  
chef_version '>= 15.0' if respond_to?(:chef_version)  
  
depends 'haproxy', '~> 8.3.0'
```

EXERCISE



Dynamic Web Load Balancer

Every time we create a web node we need to update our load balancer (myhaproxy) cookbook. That doesn't feel right!

Objective:

- ✓ Update the myhaproxy cookbook to dynamically use nodes with the company_web policy_name.
- ✓ Update the major version of the myhaproxy cookbook
- ☐ Update and push the Policyfile
- ☐ Run chef-client on the load balancer node
- ☐ Verify the load balancer node relays requests to both web nodes

GL: Ensure You are in the chef-repo



```
$ cd ~/chef-repo
```


GL: Update the Policyfile



```
$ chef update policyfiles/myhaproxy.rb
```

```
Attributes already up to date
Building policy myhaproxy
Expanded run list: recipe[myhaproxy::default]
Caching Cookbooks...
Installing myhaproxy >= 0.0.0 from path
Using      haproxy      8.3.0
Using      build-essential 8.2.1
Using      yum-epel      4.1.4
Using      seven_zip     4.2.2
Using      mingw          2.1.3
Using      windows       6.0.1

Lockfile written to
/Users/sdelfante/chef-repo/policyfiles/myhaproxy.lock.json
Policy revision id:
```

GL: Push the myhaproxy.lock.json to Chef Infra Server



```
$ chef push prod policyfiles/myhaproxy.lock.json
```

```
Uploading policy myhaproxy (08c39ccc8f) to policy group prod
Uploading policy myhaproxy (c0cb162cdb) to policy group prod
Using      build-essential 8.2.1 (4b9d5c72)
Using      haproxy          8.3.0 (1a4f7607)
Using      mingw            2.1.3 (9f5d572c)
Using      myhaproxy        1.0.0 (1a9d7377)
Using      seven_zip        4.2.2 (0e1fed3b)
Using      windows          6.0.1 (042f3380)
Using      yum-epel         4.1.4 (187c02d6)
```


GL: Converging the Load Balancer Node



```
$ knife ssh 'name:lb' -x chef -P PASSWORD 'sudo chef-client'
```

```
ec2-35-170-33-199.compute-1.amazonaws.com Starting Chef Infra Client, version 17.3.48
ec2-35-170-33-199.compute-1.amazonaws.com Using policy 'myhaproxy' at revision
'9da82055ea2c960cd680d131de8e92ba773703538575e4e429a52ca13f01fbba'
ec2-35-170-33-199.compute-1.amazonaws.com resolving cookbooks for run list:
["myhaproxy::default@1.0.0 (e9a41c2)"]
ec2-35-170-33-199.compute-1.amazonaws.com Synchronizing Cookbooks:
ec2-35-170-33-199.compute-1.amazonaws.com   - build-essential (8.2.1)
ec2-35-170-33-199.compute-1.amazonaws.com   - haproxy (8.3.0)
ec2-35-170-33-199.compute-1.amazonaws.com   - mingw (2.1.3)
ec2-35-170-33-199.compute-1.amazonaws.com   - seven_zip (4.2.2)
ec2-35-170-33-199.compute-1.amazonaws.com   - windows (6.0.1)
ec2-35-170-33-199.compute-1.amazonaws.com   - yum-epel (4.1.4)
ec2-35-170-33-199.compute-1.amazonaws.com   - myhaproxy (1.0.0)
...
...
ec2-35-170-33-199.compute-1.amazonaws.com Chef Infra Client finished, 6/26 resources
updated in 05 seconds
```

EXERCISE




Dynamic Web Load Balancer

Every time we create a web node we need to update our load balancer (myhaproxy) cookbook. That doesn't feel right!

Objective:

- ✓ Update the myhaproxy cookbook to dynamically use nodes with the company_web policy_name.
- ✓ Update the major version of the myhaproxy cookbook
- ✓ Update and push the Policyfile
- ✓ Run chef-client on the load balancer node
- ☐ Verify the load balancer node relays requests to both web nodes



← ⓘ 54.159.197.193 150% ↻


Chef Welcomes You!

PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400



← ⓘ 54.159.197.193 150% ↻

Chef Welcomes You!

PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400

← ⓘ 184.73.96.131 150% ↻


Chef Welcomes You!

PLATFORM: centos

HOSTNAME: ip-172-31-26-186

MEMORY: 604192kB

CPU Mhz: 1799.999



← ⓘ 54.159.197.193 150% ↻

Chef Welcomes You!

PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400



← ⓘ 54.159.197.193 150% ↻

Chef Welcomes You!

PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400

← ⓘ 184.73.96.131 150% ↻

Chef Welcomes You!

PLATFORM: centos


HOSTNAME: ip-172-31-26-186

MEMORY: 604192kB

CPU Mhz: 1799.999

DISCUSSION

Let's Test that Our Code Really Works



To verify that our code is working, let's remove the load balancer configuration file, forcing our code to run.

GL: Delete the haproxy.cfg File




```
$ knife ssh 'name:lb' -x chef -P PASSWORD "sudo rm  
/etc/haproxy/haproxy.cfg"
```



GL: Converge the Load Balancer



```
$ knife ssh 'name:lb' -x chef -P PASSWORD 'sudo chef-client'
```

```
ec2-34-196-50-77.compute-1.amazonaws.com Starting Chef Infra Client, version 17.3.48
...
ec2-34-196-50-77.compute-1.amazonaws.com resolving cookbooks for run list:
["myhaproxy::default@1.0.0 (e9a41c2)"]
ec2-34-196-50-77.compute-1.amazonaws.com Synchronizing Cookbooks:
...
ec2-34-196-50-77.compute-1.amazonaws.com * template[/etc/haproxy/haproxy.cfg] action
create
ec2-34-196-50-77.compute-1.amazonaws.com      - create new file /etc/haproxy/haproxy.cfg
ec2-34-196-50-77.compute-1.amazonaws.com      - update content in file
/etc/haproxy/haproxy.cfg from none to 4334de
ec2-34-196-50-77.compute-1.amazonaws.com      - suppressed sensitive resource
ec2-34-196-50-77.compute-1.amazonaws.com      - change mode from ' ' to '0644'
ec2-34-196-50-77.compute-1.amazonaws.com      - change owner from ' ' to 'haproxy'
ec2-34-196-50-77.compute-1.amazonaws.com      - change group from ' ' to 'haproxy'
...
```



  54.159.197.193 150% 


Chef Welcomes You!

PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400



  54.159.197.193 150% 

Chef Welcomes You!

PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400

  184.73.96.131 150% 


Chef Welcomes You!




PLATFORM: centos

HOSTNAME: ip-172-31-26-186

MEMORY: 604192kB

CPU Mhz: 1799.999



  54.159.197.193 

Chef Welcomes You!


PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400



  54.159.197.193 

Chef Welcomes You!

PLATFORM: windows

HOSTNAME: WIN-DQFQCUFHDCP

MEMORY: 1048176kB

CPU Mhz: 2400

  184.73.96.131 

Chef Welcomes You!


PLATFORM: centos

HOSTNAME: ip-172-31-26-186

MEMORY: 604192kB

CPU Mhz: 1799.999

EXERCISE



Dynamic Web Load Balancer


Every time we create a web node we need to update our load balancer (myhaproxy) cookbook. That doesn't feel right!

Objective:

- ✓ Update the myhaproxy cookbook to dynamically use nodes with the company_web policy_name.
- ✓ Update the major version of the myhaproxy cookbook
- ✓ Update and push the Policyfile
- ✓ Run chef-client on the load balancer node
- ✓ Verify the load balancer node relays requests to both web nodes

DISCUSSION

Review Questions



1. What is the great advantage of using the following dynamic search in the load balancer's default.rb?

```
web_nodes = search('node','policy_name:company_web')
```


```
server_array = []
```

```
web_nodes.each do |one_node|  
  one_server = "#{one_node['cloud']['public_hostname']}  
#{one_node['cloud']['public_ipv4']}:80 maxconn 32"  
  server_array.push(one_server)  
end
```

```
...
```

DISCUSSION

Review Questions



2. What is the key item that tells the load balancer how to find the web servers it's supposed to balance?


```
web_nodes = search('node','policy_name:company_web')
```

```
server_array = []
```

```
web_nodes.each do |one_node|  
  one_server = "#{one_node['cloud']['public_hostname']}  
#{one_node['cloud']['public_ipv4']}:80 maxconn 32"  
  server_array.push(one_server)  
end
```

```
...
```

DISCUSSION



Q&A

What questions can we help you answer?



CHEF™